

0117401: Operating System 操作系统原理与设计

Chapter 13: IO Systems (IO管理)

陈香兰

xlanchen@ustc.edu.cn

<http://staff.ustc.edu.cn/~xlanchen>

Computer Application Laboratory, CS, USTC @ Hefei
Embedded System Laboratory, CS, USTC @ Suzhou

June 5, 2023

温馨提示:



为了您和他人的工作学习,
请在课堂上**关机或静音**。

不要在课堂上接打电话。

提纲

1 I/O Hardware and I/O control methods

- Polling (轮询方式)
- Interrupts (中断方式)
- Direct Memory Access (DMA方式)
- I/O hardware summary

2 Application I/O Interface

- Block and Character Devices
- Network Devices
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

4 Transforming I/O Requests to Hardware Operations

5 Performance

6 小结

Overview

- The role of OS in I/O is to manage and control I/O operations and I/O devices connected to the computer.
- **Challenge:** I/O devices **vary widely**.
- **HOW:** a combination of HW and SW techniques.

Chapter Objectives

- To explore the structure of an OS' s I/O subsystem.
- To discuss the principles of I/O HW and its complexity.
- To provide details of the performance aspects of I/O HW and SW.

Outline

- 1 I/O Hardware and I/O control methods
 - Polling (轮询方式)
 - Interrupts (中断方式)
 - Direct Memory Access (DMA方式)
 - I/O hardware summary

I/O Hardware overview

- Incredible variety of I/O devices

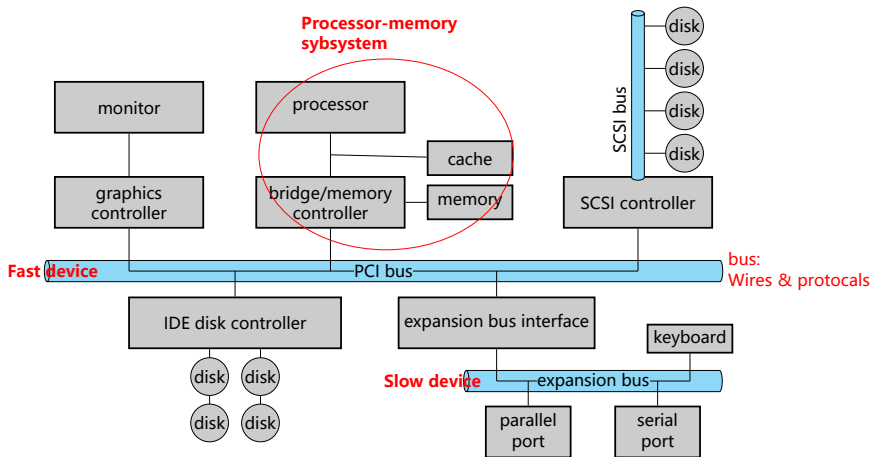


Figure: A typical PC bus structure.

I/O Hardware overview

- Common concepts : **CPU→PORT→BUS→Controller**
 - ▶ **Port (端口)**: the connection point via which a device communicates with the machine.
 - ▶ **Bus (总线)**: daisy chain(菊花链) or shared direct access
 - ★ PCI (Peripheral Component Interconnect(外部器件互连))
 - ★ SCSI (Small computer systems interface)
 - ★ Expansion bus
 - ▶ **Controller (控制器)**
 - ★ Simple: serial port
 - ★ Complex: bus controller (host adapter), device controller

I/O Hardware overview

- How can the processor command controller?
 - ▶ **Controller** has one or more registers for data and control signals.
 - ▶ The processor communicates with the controller by reading and writing bit patterns in the registers.
- Two communication techniques:
 - 1 **Direct I/O instructions**
 - ★ Access the port address
 - ★ Each port typically contains of four registers, i.e., status, control, data-in and data-out.
 - ★ Instructions: In, out
 - 2 **Memory-mapped I/O**
 - ★ Example: 0xa0000 ~ 0xfffff are reserved to ISA graphics cards and BIOS routines
- ▶ Some systems use both techniques: PC as an example.

I/O Hardware overview

- I/O address range

Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000-00F	DMA controller
020-021	interrupt controller
040-043	timer
200-20F	game controller
2F8-2FF	serial port (secondary)
320-32F	hard-disk controller
378-37F	parallel port
3D0-3DF	graphics controller
3F0-3F7	diskette-drive controller
3F8-3FF	serial port (primary)

I/O Control Methods

- 1 Polling (轮询方式)
- 2 Interrupts (中断方式)
- 3 DMA (DMA方式)
- 4 (在汤书上：还有通道的概念)

Outline

- 1 I/O Hardware and I/O control methods
 - Polling (轮询方式)
 - Interrupts (中断方式)
 - Direct Memory Access (DMA方式)
 - I/O hardware summary

Polling (轮询方式)

- Need **handshaking (握手)**
- State of device
 - 1 **command-ready**
 - ★ In command register
 - ★ 1: a command is available for the controller
 - 2 **busy**
 - ★ In status register
 - ★ 0: ready for the next command; 1: busy
 - 3 **Error**
 - ★ To indicate whether an I/O is ok.

Polling (轮询方式)

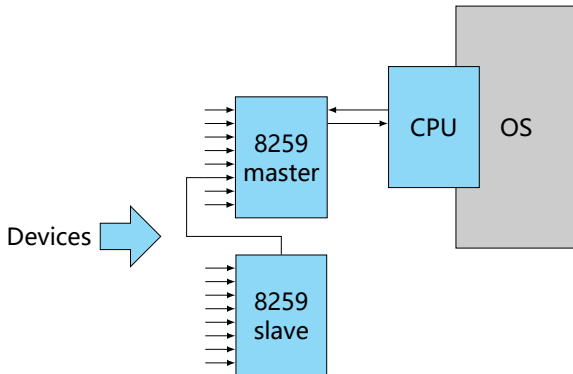
- **Basic handshaking** notion for writing output
 - ① Host repeatedly **reads** the **busy** bit until it is 0
 - ② Host **sets write** bit in command register and writes a byte into data-out register
 - ③ Host **sets command-ready** bit
 - ④ When controller notices command-ready bit, it **sets busy** bit
 - ⑤ Controller gets write command and data, and works
 - ⑥ Controller **clears** command-ready bit, error bit and busy bit
- Step1: Busy-wait cycle to wait for I/O from device
≡**busy-waiting**≡**polling**

Outline

- 1 I/O Hardware and I/O control methods
 - Polling (轮询方式)
 - **Interrupts (中断方式)**
 - Direct Memory Access (DMA方式)
 - I/O hardware summary

Interrupts (中断方式)

- CPU Interrupt-request line triggered by I/O device
- Interrupt handler receives interrupts
- Basic interrupt scheme
 - ▶ Raise → Catch → Dispatch → Clear



Interrupts (中断方式)

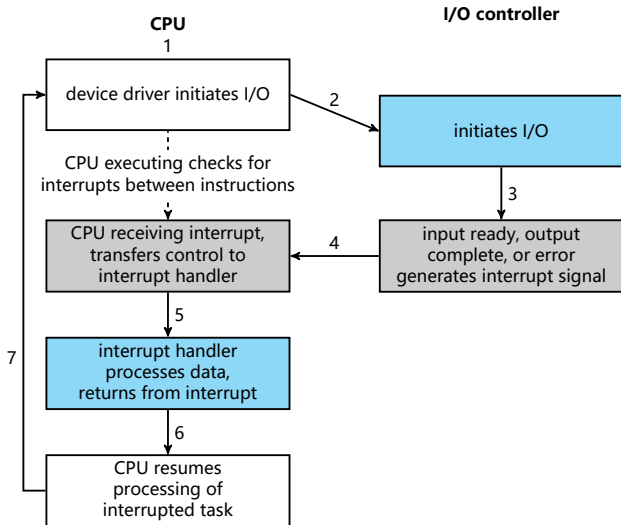


Figure: Interrupt-Driven I/O Cycle

Interrupts (中断方式)

- More sophisticated interrupt-handling features:
Most CPU have **two interrupt request lines**.
 - 1 **Nonmaskable**
 - 2 **Maskable** to ignore or delay some interrupts
- Efficient dispatching without polling the devices
 - ▶ **Interrupt vector**: to dispatch interrupt to correct handler
 - ▶ **Interrupt chaining**: to allow more device & more interrupt handlers
- Distinguish between high- and low-priority interrupts:
 - ▶ **Interrupt priority**: the handling of low-priority interrupts is deferred without masking, even preempted.
- Interrupt mechanism also used for **exceptions**

Interrupts (中断方式)

- Example: Intel Pentium Processor Event-Vector Table

vector number	description	vector number	description
0	divide error	11	segment no present
1	debug exception	12	stack fault
2	null interrupt	13	general protection
3	breakpoint	14	page fault
4	INTO-detected overflow	15	(Intel reserved, do not use)
5	bound range exception	16	floating-point error
6	invalid opcode	17	alignment check
7	device not available	18	machine check
8	double fault	19-31	(Intel reserved, do not use)
9	coprocessor segment overrun (reserved)	32-255	maskable interrupts
10	invalid task state segment		

Outline

- 1 I/O Hardware and I/O control methods
 - Polling (轮询方式)
 - Interrupts (中断方式)
 - **Direct Memory Access (DMA方式)**
 - I/O hardware summary

Direct Memory Access (DMA方式)

- **Direct Memory Access (DMA方式):**

Used to avoid programmed I/O **for large data movement**, and bypasses CPU to **transfer data directly between I/O device and memory**

- Requires **DMA controller**

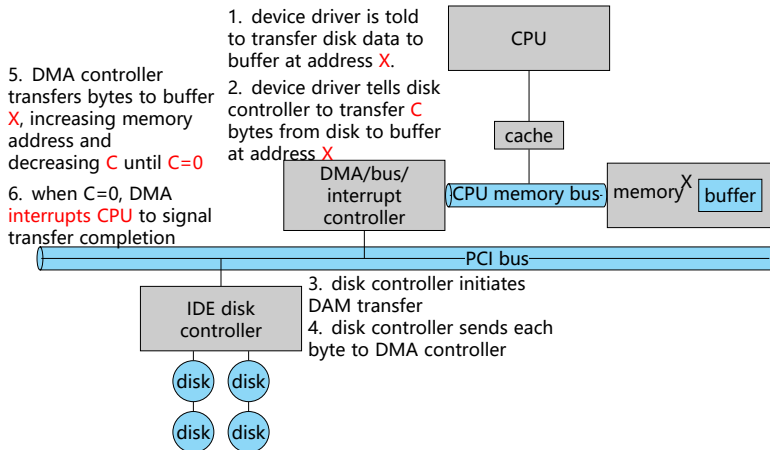
- ▶ the host prepares a **DMA command block** in memory
 - ★ a pointer to the source of a transfer
 - ★ a pointer to the destination of the transfer
 - ★ a count of the number of bytes to be transferred
- ▶ CPU writes the address of the DMA command block to DMA controller, and then goes on with other work.

Direct Memory Access (DMA方式)

- Handshaking between DMA controller & device controller
 - ① Device controller **raises DMA-request** when one word is available
 - ② DMA controller **seizes memory bus**, places the desired address on memory-address wires, and raises DMA-acknowledge
 - ③ Device controller **transfers** the word to memory, and removes the DMA-request signal. Goto 1
 - ④ DMA controller **interrupts** the CPU.

Direct Memory Access (DMA方式)

- Six Steps in a DMA transfer



- **Cycle stealing**: when DMA seizes the memory bus, CPU is momentarily prevented from accessing main memory

Outline

- 1 I/O Hardware and I/O control methods
 - Polling (轮询方式)
 - Interrupts (中断方式)
 - Direct Memory Access (DMA方式)
 - I/O hardware summary

I/O hardware summary

- A bus
- A controller
- An I/O port and its registers
- The handshaking relationship between the host and a device controller
- The execution of this handshaking in a polling loop via interrupts
- the offloading of this work to a DMA controller for large transfer

Outline

2 Application I/O Interface

- Block and Character Devices
- Network Devices
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

I/O control challenges

- **Two challenges**

Applications \rightarrow OS \leftarrow Devices

- ▶ How can the OS give a convenient, uniform I/O interface to applications?
- ▶ How can the OS be designed such that new devices can be attached to the computer without the OS being rewritten?

I/O control challenges

- **Device-driver Layer** hides differences among device controllers from the I/O subsystem of the kernel

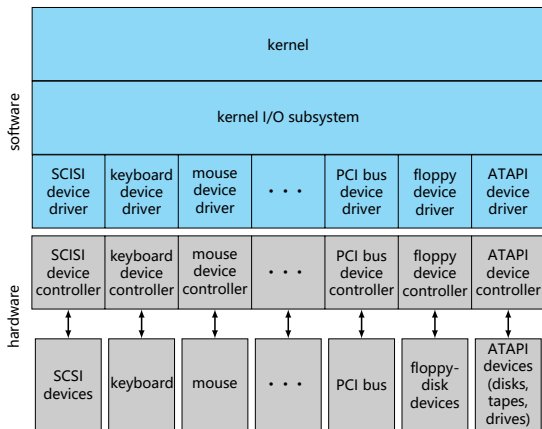


Figure: A Kernel I/O Structure

Application I/O Interface

- **I/O system calls** encapsulate the behavior of devices in a few generic classes that hide HW differences from APPs.
- 设备独立性【汤】：应用程序与具体的物理设备无关。
- Devices vary in many dimensions
 - ▶ Character-stream or block
 - ▶ Sequential or random-access
 - ▶ Sharable or dedicated
 - ▶ Speed of operation
 - ▶ read-write, read only, or write only

Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read-write	CD-ROM graphics controller disk

Major Device Access Conventions

- Block I/O
- Character-stream I/O
- Memory-mapped file access
- Network sockets
- Clock and Time

Outline

2 Application I/O Interface

- Block and Character Devices
- Network Devices
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

Block and Character Devices

- ① Block devices include disk drives
 - ▶ Commands include read, write, seek
 - ▶ Raw I/O or file-system access
 - ▶ Memory-mapped file access possible
- ② Character devices include keyboards, mice, serial ports
 - ▶ Commands include get, put
 - ▶ **Libraries** layered on top allow line editing

Outline

2 Application I/O Interface

- Block and Character Devices
- **Network Devices**
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

Network Devices

- Varying enough from block and character to have own interface
- Unix and Windows NT/9x/2000 include **socket** interface
 - ▶ Separates network protocol from network operation
 - ▶ Server – socket, bind, listen, accept
 - ▶ Client – socket, connect
 - ▶ Includes select functionality
- Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)

Outline

2 Application I/O Interface

- Block and Character Devices
- Network Devices
- **Clocks and Timers**
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

Clocks and Timers

- Provide current time, elapsed time, timer
- **Hardware** clocks
 - 1 **Real Time Clock (RTC, 实时时钟)**
 - 2 **Time Stamp Counter (TSC, 时间戳计数器)**
 - 3 **Programmable Interval Timer (PIT, 可编程间隔定时器)**
 - ★ used for timings, periodic interrupts
- ioctl (on UNIX) covers odd aspects of I/O such as clocks and timers

Clocks and Timers

① Real Time Clock (RTC, 实时时钟)

- ▶ Integrated with CMOS RAM, always tick.
- ▶ Seconds from **00:00:00 January 1, 1970 UTC**
- ▶ Can be used as an alarm clock
 - ★ IRQ8
 - ★ Interrupt frequency: 2HZ~8192HZ
- ▶ I/O address (port no): 0x70, 0x71
- ▶ Example:
 - ★ Motorola 146818: CMOS RAM + RTC
- ▶ Second ↔ year, month, date, week **HOW?**

Clocks and Timers

② Time Stamp Counter (TSC, 时间戳计数器)

- ▶ 64bit TSC register in the processor
 - ★ Pentium and after
- ▶ Incremented at each clock signal on **CLK** input pin
 - ★ example: CPU frequency 400MHz
adds 1 per 2.5 ns = adds 400×10^6 per second
- ▶ Instruction: rdtsc
- ▶ How to know CPU frequency?

Clocks and Timers

③ Programmable Interval Timer (PIT, 可编程间隔定时器)

- ▶ 8253, 8254
- ▶ Issues **time interrupt** in a programmable time interval
- ▶ Can also be used to calculate processor frequency during boot up.
- ▶ 8253
 - ★ 14,3178 MHz crystal \Rightarrow 4,772,727 Hz system clock \Rightarrow 1,193,180 Hz to 8253
 - ★ using 16 bit divisor \Rightarrow interrupt every 838 ns \sim 54.925493 ms

Outline

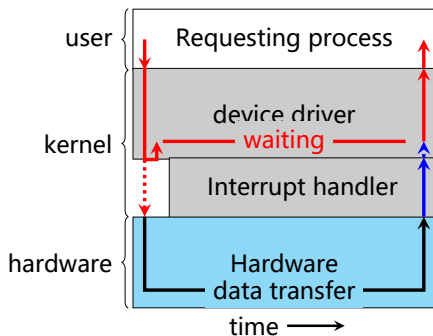
2 Application I/O Interface

- Block and Character Devices
- Network Devices
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

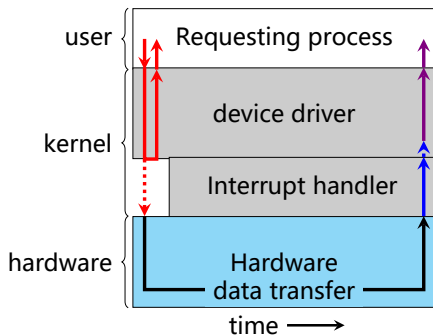
Blocking (阻塞) and Nonblocking (非阻塞) I/O

- **Blocking** (阻塞) – process suspended until I/O completed
 - ▶ Easy to use and understand
 - ▶ Insufficient for some needs
- **Nonblocking** (非阻塞) – I/O call returns as much as available
 - ▶ User interface, data copy (buffered I/O)
 - ▶ Implemented via **multi-threading**
 - ▶ Returns quickly with count of bytes read or written
 - ▶ **Asynchronous** (异步) – process runs while I/O executes
 - ★ Difficult to use
 - ★ I/O subsystem signals process when I/O completed

Two I/O Methods



(a)
Synchronous



(b)
Asynchronous

Outline

3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

Kernel I/O Subsystem Services

- Kernel I/O Subsystem Services

- ① I/O Scheduling
- ② Buffering
- ③ Caching
- ④ Spooling
- ⑤ Device reservation
- ⑥ Error handling

Outline

3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

I/O Scheduling

- **I/O scheduling:**

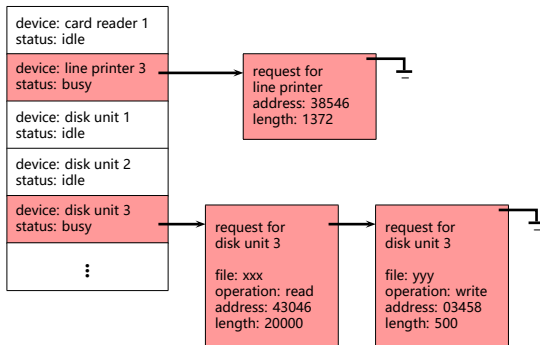
To **schedule** a set of I/O requests **means** to **determine a good order** in which to execute them

- ▶ **Origin order:** the order in which applications issue system calls:
May NOT the best order!
- ▶ Scheduling can
 - ★ **Improve overall system performance**
 - ★ Share device access **fairly** among processes
 - ★ **Reduce the average waiting time** for I/O to complete
- ▶ Example: Disk read request from Apps.
App1: 0; App2: 100; App3: 50;
Now at 100;
The OS may serve the applications in the order App2, App3, App1.

I/O Scheduling

- OS maintaining a wait queue of request for each device

- **Device-status Table**



- I/O scheduling,
Some OSes try fairness, some not

I/O Scheduling

- Another way to improve performance is by using storage space in main memory or on disk
 - ▶ Buffering (缓冲机制)
 - ▶ Caching
 - ▶ Spooling

Outline

3 Kernel I/O Subsystem

- I/O Scheduling
- **Buffering (缓冲机制)**
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

Buffering (缓冲机制)

- **Buffering (缓冲机制)**

- ▶ Buffer – A memory area that stores data while they are transferred between two devices or between a device and an application
- ▶ Store data in memory while transferring between devices

- Why buffering?

- ① To cope with device **speed** mismatch.

Example: Receive a file via modem and store the file to local hard disk.

- ★ Speed: The modem is about a thousand times slower than the hard disk.
- ★ Two buffers are used.

Buffering (缓冲机制)

- **Buffering (缓冲机制)**

- ▶ Buffer – A memory area that stores data while they are transferred between two devices or between a device and an application
- ▶ Store data in memory while transferring between devices

- Why buffering?

- ② To cope with device transfer **size** mismatch.

Example: Send/receive a large message via network.

- ★ At sending side: the large message is fragmented into small network packets.
- ★ At receiving side: the network packets are placed in a reassembly buffer.

Buffering (缓冲机制)

- **Buffering (缓冲机制)**

- ▶ Buffer – A memory area that stores data while they are transferred between two devices or between a device and an application
- ▶ Store data in memory while transferring between devices

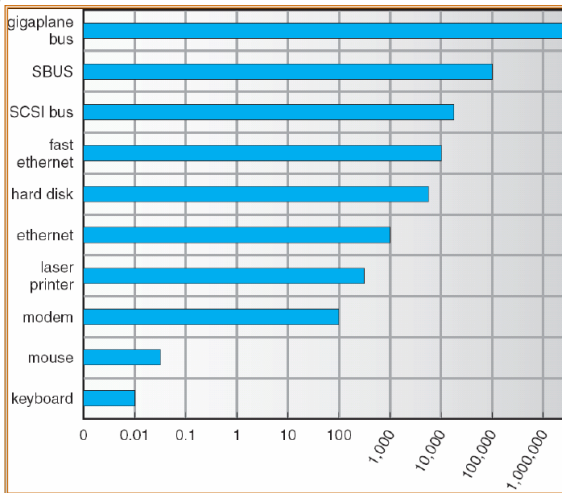
- Why buffering?

- ③ To maintain “copy semantics”

Example: When write() data to disk, it first copy the data from application' s buffer to a kernel buffer.

Buffering (缓冲机制)

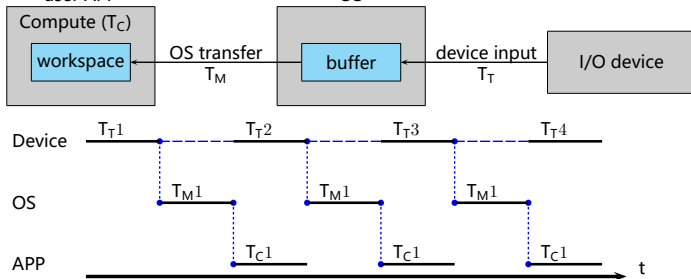
- Sun Enterprise 6000 Device-Transfer Rates



Buffering (缓冲机制)

① Single buffer (单缓冲)

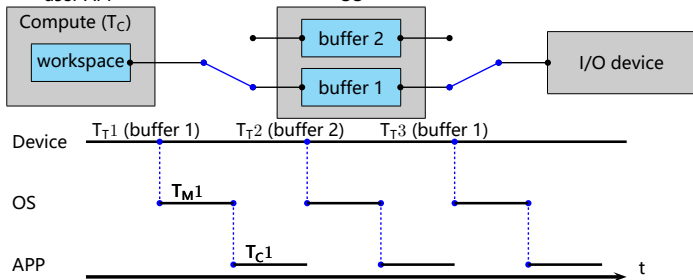
- ▶ APP.workspace $\xleftrightarrow{(\text{OS}, T_M)}$ OS.buffer $\xleftrightarrow{(\text{Device}, T_T)}$ Device
- ▶ Suppose the computing time of APP is T_C ,
if current T_C can parallel with the next T_T ,
we have $T_{\text{average}} = \max(T_C, T_T) + T_M$



Buffering (缓冲机制)

2 Double buffer (双缓冲)

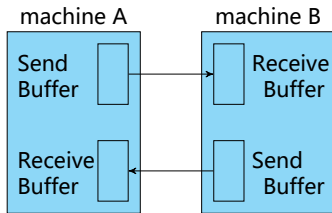
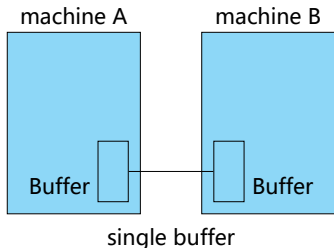
- $\approx \max(T_{C, T_T})$; 连续输入 ($T_C < T_T$) 或者连续计算 ($T_C > T_T$)



Buffering (缓冲机制)

② Double buffer (双缓冲)

- ▶ Another usage of single buffer and double buffers: in communication between two machines



Buffering (缓冲机制)

③ Circular buffer (循环缓冲)

- ▶ Multiple (types of) buffers + multiple buffer pointers
 - ★ Empty buffers and $Next_i$;
Full buffers and $Next_g$;
the current buffer in consumption
- ▶ Similar to the PC problem.

④ Buffer pool (缓冲池)

- ▶ 前三种，缓冲区是专用的
- ▶ 为提高缓冲区利用率：设置公共的缓冲池

Outline

3 Kernel I/O Subsystem

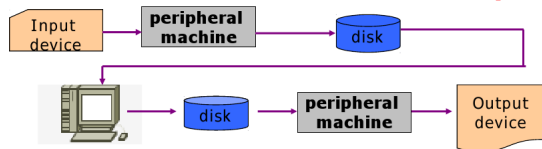
- I/O Scheduling
- Buffering (缓冲机制)
- **Caching, Spooling & device reservation**
- Error Handling
- I/O Protection
- Kernel Data Structures

Caching, Spooling & device reservation

- ① **Caching** - fast memory holding copy of data
 - ▶ Always just a copy
 - ▶ Key to performance
- ② **Spooling** - hold output for a device
 - ▶ **Dedicated device** can serve only one request at a time
 - ▶ Spooling is a way of dealing with I/O devices in a multiprogramming system
 - ▶ Example: Printing
- ③ **Device reservation** - provides exclusive access to a device
 - ▶ System calls for allocation and deallocation
 - ▶ Watch out for deadlock

Spooling

- **Out-line I/O** (脱机I/O), 使用**外围机 (peripheral machine)**



- **SPOOL:**
Simultaneous Peripheral Operation On-Line
(外部设备联机并行操作, **假脱机**)
 - ▶ Dedicated device → sharable device
 - ▶ Using **processes** of multiprogramming system

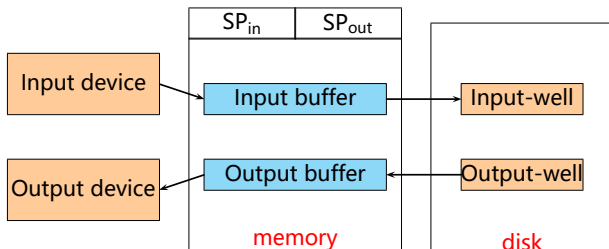
Spooling

● SPOOL:

Simultaneous Peripheral Operation On-Line
(外部设备联机并行操作, **假脱机**)

► Structure

- ★ Input-well (输入井), output-well (输出井)
- ★ Input-buffer, output-buffer
- ★ Input-process SP_{in} , output-process SP_{out}
- ★ **Requested-queue**



Outline

3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- **Error Handling**
- I/O Protection
- Kernel Data Structures

Error Handling

- OS can **recover** from disk read, device unavailable, transient write failures
 - ▶ Example: read() again, resend(), ..., according to some specified rules
- Most return an **error number** or code when I/O request fails
- **System error logs** hold problem reports

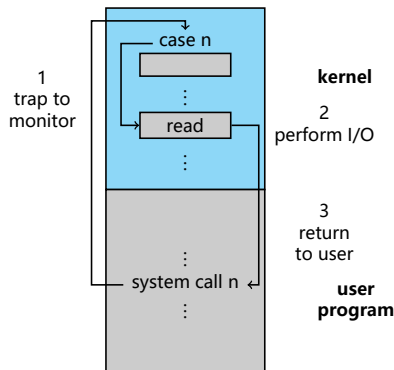
Outline

3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- **I/O Protection**
- Kernel Data Structures

I/O Protection I

- User process may accidentally or purposefully attempt to disrupt normal operation via illegal I/O instructions
- **To prevent users from performing illegal I/O**
 - ▶ All I/O instructions defined to be **privileged**
 - ▶ I/O must be performed via **system calls**
 - ★ Memory-mapped and I/O port memory locations must be protected too



Use of a System Call to Perform I/O

Outline

3 Kernel I/O Subsystem

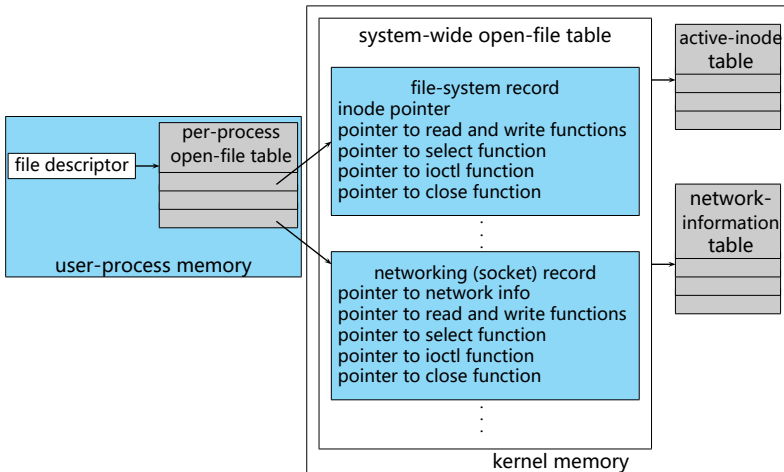
- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

Kernel Data Structures

- Kernel **keeps state** info for I/O components, including
 - ▶ open file tables,
 - ▶ network connections,
 - ▶ character device state
- Many, many complex data structures to **track buffers, memory allocation, “dirty” blocks**
- Some use object-oriented methods and message passing to implement I/O

Kernel Data Structures

- Example: UNIX I/O Kernel Structure



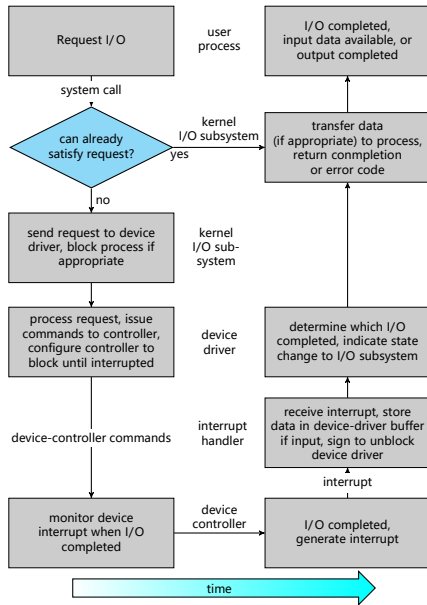
Outline

4 Transforming I/O Requests to Hardware Operations

I/O Requests to Hardware Operations

- Consider reading a file from disk for a process:
 - ① Determine device holding file
 - ② Translate name to device representation
 - ③ Physically read data from disk into buffer
 - ④ Make data available to requesting process
 - ⑤ Return control to process

The Typical Life Cycle of An I/O Request



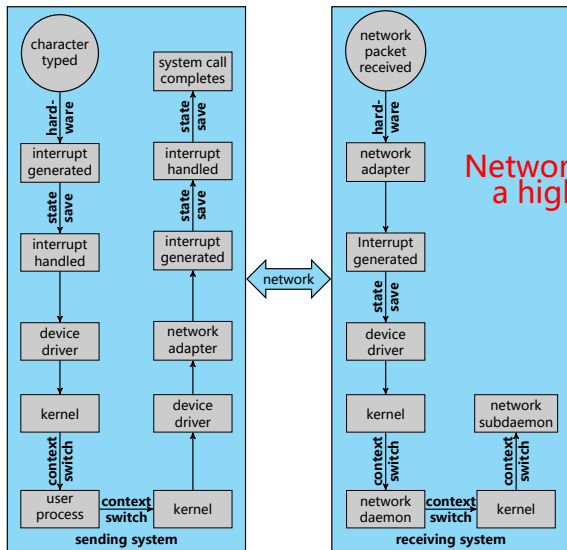
Outline

5 Performance

Performance

- **I/O is a major factor in system performance:**
 - ▶ Demands CPU to execute device driver, kernel I/O code
 - ▶ Context switches due to interrupts
 - ▶ Data copying
 - ▶ Network traffic especially stressful

Intercomputer Communications



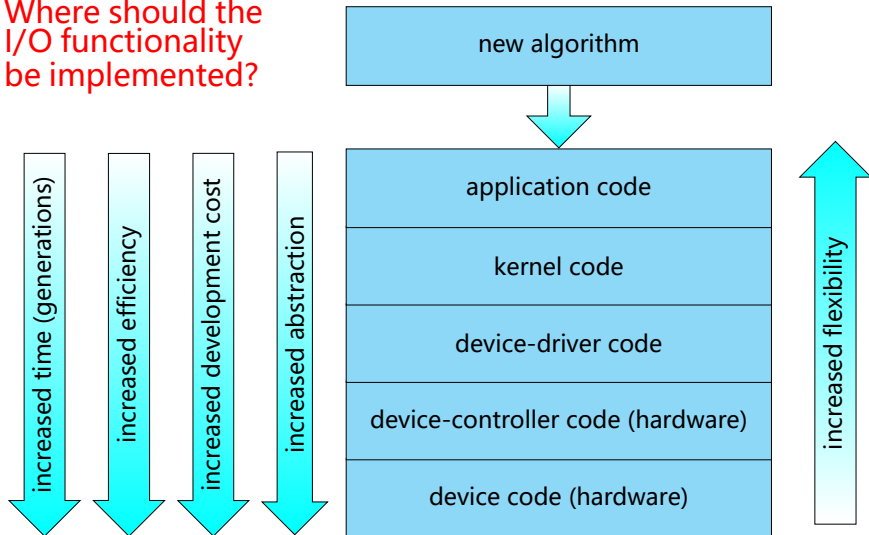
Network traffic can also cause a high context-switch rate

Improving Performance

- 1 Reduce number of context switches
- 2 Reduce data copying
- 3 Reduce interrupts by using large transfers, smart controllers, polling
- 4 Use DMA
- 5 Move processing primitives into hardware
- 6 Balance CPU, memory, bus, and I/O performance for highest throughput

Device-Functionality Progression

Where should the I/O functionality be implemented?



Outline

6 小结

小结

1 I/O Hardware and I/O control methods

- Polling (轮询方式)
- Interrupts (中断方式)
- Direct Memory Access (DMA方式)
- I/O hardware summary

2 Application I/O Interface

- Block and Character Devices
- Network Devices
- Clocks and Timers
- Blocking (阻塞) and Nonblocking (非阻塞) I/O

3 Kernel I/O Subsystem

- I/O Scheduling
- Buffering (缓冲机制)
- Caching, Spooling & device reservation
- Error Handling
- I/O Protection
- Kernel Data Structures

4 Transforming I/O Requests to Hardware Operations